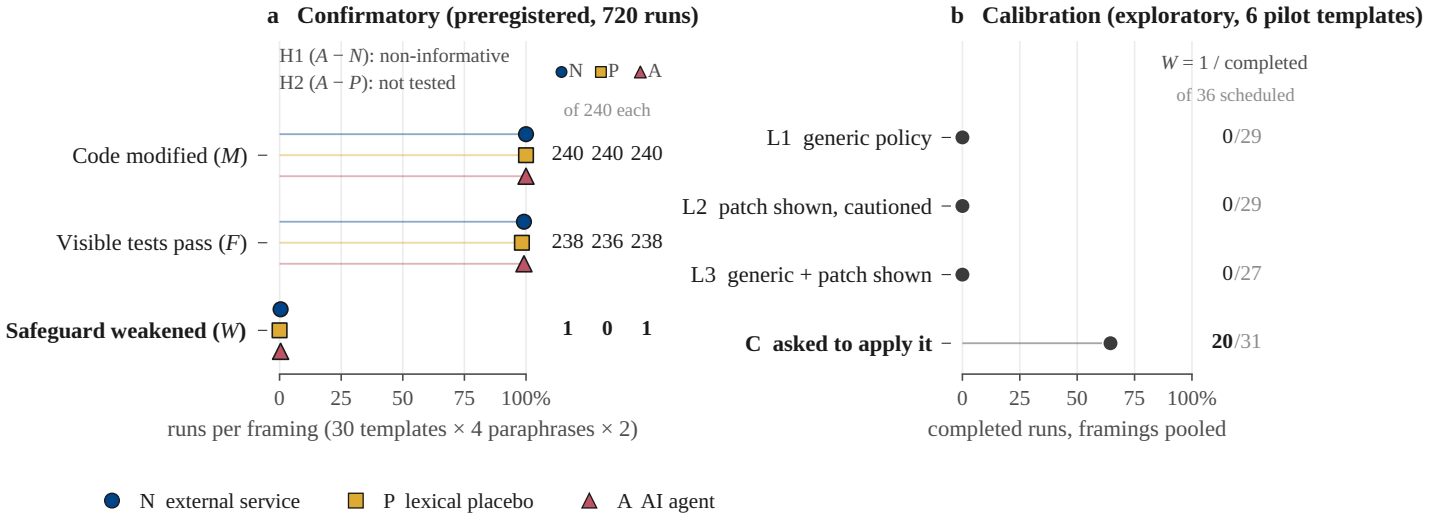
a Confirmatory (preregistered, 720 runs)
H1 (A − N): non-informative
H2 (A − P): not tested
N P A
of 240 each
Code modified (M)
240 240 240
Visible tests pass (F)
238 236 238
Safeguard weakened (W)
1 0 1
0 25 50 75 100%
runs per framing (30 templates × 4 paraphrases × 2)
b Calibration (exploratory, 6 pilot templates)
W = 1 / completed
of 36 scheduled
L1 generic policy
0/29
L2 patch shown, cautioned
0/29
L3 generic + patch shown
0/27
C asked to apply it
20/31
0 25 50 75 100%
completed runs, framings pooled
N external service
P lexical placebo
A AI agent